

The Clay Navier–Stokes Problem: Formulation and Solution

Claes Johnson¹

¹KTH Royal Institute of Technology, SE-100 44 Stockholm, Sweden ,
claesjohnson@gmail.com

Abstract

The official formulation of the Clay Millennium Problem on the Navier–Stokes equations asks whether smooth initial data give a solution that stays smooth for all time, or whether the velocity can blow up in finite time. The alternative is exhaustive only if those are the two possibilities. They are not. The solutions of physical interest are *turbulent*: not smooth, because velocity gradients grow as $\nu^{-1/2}$ and are enormous at the Reynolds numbers of any real flow, and not singular, because the velocity itself stays bounded. A dichotomy that omits the only case that matters cannot be answered in a way that says anything about fluids, and the case it does cover — persistent smoothness — is either trivial or false depending on how it is read.

This article argues that the problem as posed is mis-posed, and states what should replace it. The replacement is Hadamard’s criterion: not global smoothness, but *wellposedness*. A solution concept earns physical meaning when the quantities computed from it are stable under perturbation. For turbulent solutions those quantities are not pointwise velocities, which are indeterminate, but mean-value outputs such as drag and lift, which are not. We set out the distinction, the evidence that turbulent solutions are wellposed in outputs while smooth solutions are not, and the computational record behind it: a Galerkin method whose only turbulence model is its own residual, error control by duality with computed stability factors, and mean drag and lift converging to a few per cent on meshes on which the velocity field does not converge and never will. The reformulated problem has a solution, and the solution is that turbulence is what Navier–Stokes has to say.

1 The official formulation

The Clay Mathematics Institute problem, as set out by Fefferman [1], concerns the incompressible Navier–Stokes equations

$$\dot{u} + (u \cdot \nabla)u - \nu \Delta u + \nabla p = f, \quad \nabla \cdot u = 0, \quad (1)$$

on $\mathbb{R}^3$ or a periodic box, with $u(\cdot, 0) = u^0$ smooth and divergence free, and with the viscosity ν normalised to one. Four statements are offered, of which a proof of any one settles the prize: existence and smoothness on $\mathbb{R}^3$; the same with periodic data; or a counterexample to either — a solution which is initially smooth and whose velocity becomes unbounded in finite time.

The structure of the question is therefore a dichotomy. Either the solution is smooth for all time, or it is not, and the only stated way for it not to be is blowup. *Tertium non datur*.

2 Why the dichotomy is empty

The physical solution is neither. A fluid at a Reynolds number of interest — 10^6 and upward for an aircraft, a ship, a weather system, a pipe — is turbulent. Turbulent flow is not smooth in any useful sense: the velocity varies on a length scale that shrinks with viscosity, and the gradients scale as

$$|\nabla u| \sim \nu^{-1/2}, \quad (2)$$

so that at $\nu = 10^{-6}$ they are of order 10^3 times the velocity scale, and at the viscosity of air over a wing, larger still. Nor is it singular: the velocity stays bounded, the kinetic energy stays finite, and nothing goes to infinity anywhere. Kolmogorov’s scaling, in the form Onsager gave it, puts the velocity field at Hölder continuity $1/3$ — continuous, with unbounded derivatives in the limit. That is precisely the case the official alternative does not contain.

The smooth case is trivial or false. If smoothness is to be read as “the solution exists and is smooth for a short time”, the statement is classical and was not worth a prize. If it is to be read as “the solution stays smooth for all time”, then for the flows of interest it is false, and not because of blowup: a smooth solution of (1) at small viscosity is *unstable*. Perturb potential flow around a bluff body and it does not stay potential flow; it develops rotational separation and becomes turbulent. The smooth solution exists as a formal object and is not observed, because nothing keeps a flow in it. So the first horn of the dichotomy describes solutions that are mathematically available and physically absent, and the second describes solutions that do not occur.

The normalisation is legitimate; what is wrong is the quantifier. It might be thought that setting $\nu = 1$ removes the phenomenon, since all the structure in (2) is structure in ν . Taken as an objection to the normalisation itself, that is wrong. Navier–Stokes rescales as

$$v(x, t) = \alpha u(\beta x, \alpha \beta t), \quad \nu' = \nu \beta / \alpha,$$

so any viscosity can be set to one by rescaling the data: small ν with data of size one is the *same problem* as $\nu = 1$ with large data. Since Fefferman quantifies over all Schwartz-class data, arbitrarily high Reynolds number is inside the hypotheses, and the family of problems is the right one.

The defect is what is asked of that family. Fefferman asks, *for each* datum separately, whether a global smooth solution exists. Turbulence is not a property of any one member; it is a property of the family as the Reynolds number grows, and a question about it must be asked *uniformly*. Put the family in its natural coordinates — data of size one, $\nu > 0$ with no positive lower bound — and the distinction is immediate:

quantity	bounded for each fixed $\nu > 0$?	bounded uniformly as $\nu \rightarrow 0$?
velocity $\ u\ $	yes	yes (energy)
gradient $\ \nabla u\ $	yes	no — grows as $\nu^{-1/2}$
mean-value output $M(u)$	yes	yes — and this is the content

Read down the first column and every answer is yes, which is Fefferman’s question and which settles nothing. Read down the second and the three quantities separate, which is the question about fluids. A theorem of the form “for each ν the solution is smooth, with derivatives bounded by an iterated exponential” is true in the first column and empty in the second, since the bound degenerates precisely in the limit at issue. That is the vacuity, and it survives the scaling argument rather than depending on it.

Direct simulation settles the matter by exhibition. The argument need not stay abstract, because solutions of the kind at issue are computed routinely. Direct numerical simulation resolves the flow down to the Kolmogorov cutoff $\eta \sim \nu^{3/4}$, which requires $L/\eta \sim \text{Re}^{3/4}$ points per direction and $\text{Re}^{9/4}$ in all — a thousand points a side at $\text{Re} = 10^4$, which is ordinary. The field so computed is *smooth by construction*: it is fully resolved, every derivative is finite, nothing is singular anywhere, and no structure is left below the mesh.

It is also, on inspection, turbulent. One can look at it.

So here are solutions which satisfy the first horn of the dichotomy exactly — smooth, existing for all time, no blowup — and which are the very phenomenon the problem is supposed to be about. If (A) were proved tomorrow it would be affirmed by these fields, and would say nothing whatever about them that distinguishes them from laminar flow. The predicate does not separate the cases it is being asked to separate. That is not an objection to the proof techniques; it is an objection to the question, and direct simulation makes it not an argument but an observation.

The real gap, and the question it should have become. It is worth being exact about what *is* open here, because there is a genuine gap and the Clay problem was posed in its vicinity. That turbulent flow has a smallest scale set by the viscosity is doubted by almost nobody, and on very substantial grounds: dimensional analysis, the observed independence of the dissipation rate from ν , and a century of measurement. But it is not a theorem about (1). Kolmogorov’s $\eta \sim \nu^{3/4}$ rests on local isotropy and universality assumed, not derived; the zeroth law is an empirical regularity; and that Navier–Stokes solutions converge as $\nu \rightarrow 0$ to anything in particular is unproved. A mathematical question is genuinely wanted here, and that is the impulse behind the prize.

What is proved lies on the Euler side and is worth setting against the dichotomy. Onsager’s threshold is settled in both directions: above Hölder $1/3$ the energy is conserved [3], and below it dissipating solutions exist and have been constructed [4]. So the third case — bounded, not smooth, dissipating — is not a conjecture about what solutions might do. It is a category with occupants.

The question the gap asks is therefore: *does the cascade terminate at a scale set by the viscosity, and does the field stay bounded while its gradients do not?* That is a question about the structure of the solution, and either answer is informative. Fefferman’s formulation asks instead whether the velocity stays finite. The two are not the same question, and the second cannot express the first: a solution whose cascade fails to terminate is non-smooth without being singular, and the dichotomy has no name for it. The inspiration was sound and the formalisation lost it.

The quantitative results confirm rather than relieve this. Work in the analytic tradition has produced conditional regularity bounds of the form: if the velocity is bounded by A , then the derivatives are bounded by an expression exponential in A iterated several times. As a theorem this is real. As a statement about a fluid it is empty: a bound of the size $\exp \exp \exp A$ is not a bound that any flow can be said to respect, and no measurement could distinguish it from no bound at all. The objection sometimes made to the problem’s framing — that the selected question has little physical content, and reads as one posed from harmonic analysis rather than from continuum mechanics — is, on this evidence, well taken.

Non-smooth and non-singular, quantitatively. The two halves of the third case are one scaling seen from either end. Dissipation at a rate independent of the viscosity, $\varepsilon \sim \nu |\nabla u|^2$ held fixed as $\nu \rightarrow 0$, gives (2); and a structure across which the velocity jumps by $O(1)$ therefore has

thickness

$$\delta \sim |\nabla u|^{-1} \sim \nu^{1/2}. \quad (3)$$

At $\nu = 10^{-6}$ that is a thousandth of the macroscopic scale; at 10^{-12} , a millionth. As $\nu \rightarrow 0$ the structure thins without limit while the velocity jump across it stays bounded: the gradient diverges, so the solution is not smooth, and the velocity does not, so it is not singular. That is the whole of the third case in one line, and (3) is the reason no amount of regularity theory reaches it.

It also disposes of a tempting misreading. A singular structure produced by letting $\nu \rightarrow 0$ is not a counterexample to (1) at any fixed viscosity, because by (3) its thickness at that viscosity is $\nu^{1/2}$ and finite, and the solution there is bounded. Such a structure is singular only in the limit, where the equation being solved is Euler’s and not Navier–Stokes’s. The Clay problem is posed at fixed positive viscosity. A construction whose singularity shrinks to nothing as the viscosity is removed has answered the Euler question — itself open, and interesting — and not the one asked. The same test applies to any claimed blowup: does the singular structure survive at fixed $\nu > 0$ with data satisfying the stated decay, or does it exist only in a limit, or only with a forcing term the formulation excludes?

What the problem advertises, and what it asks. The Clay Institute’s own description of the problem opens on turbulence. Waves follow the boat across the lake, turbulent air currents follow the aircraft, and mathematicians and physicists believe an explanation of both the breeze and the turbulence is to be found in solutions of the Navier–Stokes equations; the challenge is to unlock the secrets hidden in them. That is the advertisement, and it is accurate about where the secrets are.

The formal statement that follows contains no turbulence. It asks whether the velocity stays finite. Between the description and the problem, the subject has been sorted out — not by oversight in the description, which names the phenomenon correctly, but by the formalisation, which admits no predicate that could distinguish a turbulent solution from a laminar one. The reader is promised the secret hidden in Navier–Stokes and handed a regularity question. That is the unfortunate part, and it is why this article argues for reformulation rather than for a different proof strategy.

The alternatives are not complementary, and the prize can be won without settling anything. Fefferman’s four statements do not form two pairs of negations. (A) and (B) are posed with $f \equiv 0$ and quantify over *all* admissible data: does every datum give a global smooth solution? (C) and (D) ask instead for the *existence* of admissible data *together with an admissible smooth forcing* f for which smoothness fails. So (C) is not $\neg$ (A). A blowup driven by an external force leaves (A) untouched — it does not show that a fluid tears itself into a singularity, only that something can be done to it that does.

This is not hypothetical. In September 2026 a machine-generated and Lean-verified proof of (C) was announced, constructing, for *every* positive viscosity, a flow starting from rest and driven by a smooth compactly supported force whose velocity becomes unbounded in finite time at bounded kinetic energy. Whether it stands is for the referees. What it already shows is the present point: an alternative the problem itself offers has been discharged, and the question everyone means by the Navier–Stokes problem — can the equations, unforced, produce a singularity — is exactly where it was.

There is a second thing to notice, and it belongs to this article’s subject rather than to the referees’. Such a construction gives, at each viscosity, a solution whose velocity becomes unbounded. Follow the family down as $\nu \rightarrow 0$ and the limit is singular. But that is not what happens to a fluid as the viscosity is reduced: what happens is turbulence, which by (2) and (3) has diverging

gradients and bounded velocity, and by the Onsager results above is a category with occupants. So the family constructed has no physical counterpart at either end — it is singular where fluids are turbulent — and the fact that it nonetheless answers one of the four stated alternatives is a fact about the alternatives. A formulation whose horns can be satisfied without bearing on its subject is a formulation in need of repair, and that is an observation about the statement, not about the proof.

(A) is provable, for an arbitrarily small change to the equation. There is a third way the dichotomy fails to bite, and it is the oldest. Augment the viscous stress by a term depending on a power of the local velocity gradient,

$$\dot{u} + (u \cdot \nabla)u - \nu \Delta u - \nabla \cdot (|\nabla u|^{p-2} \nabla u) + \nabla p = f, \quad (4)$$

so that the effective viscosity is $\nu + c|\nabla u|^{p-2}$ rather than ν . Then the standard machinery of monotone operators and compactness — Galerkin approximation, an energy estimate, a monotonicity argument to pass to the limit in the nonlinearity — delivers global existence and uniqueness of regular solutions. Lions [9] proved this in 1969 for

$$p \geq \frac{n+2}{2} = \frac{5}{2} \quad (n=3), \quad (5)$$

that is, for viscosity augmented by $|\nabla u|^{1/2}$; later work lowered the threshold to $p \geq 11/5$, an augmentation by $|\nabla u|^{1/5}$. The exponent is small, and it can be multiplied by a coefficient c as small as one likes without affecting the theorem: the estimate needs the term's *presence*, not its size.

So Fefferman's alternative (A) — existence and smoothness — is not hard in itself. It is provable by techniques that were standard before the problem was posed, for an equation differing from (1) by a term with an arbitrarily small coefficient and an arbitrarily small power. What the prize turns on is the exact value $p = 2$ and not a hair more, which is a statement about a limiting case of a family whose every other member is settled.

Two things follow. The first is that the difficulty is concentrated at a point of no physical significance: nothing measurable distinguishes $c = 0$ from $c = 10^{-10}$, and the fluid does not know which equation it is supposed to satisfy. The second is sharper, and returns to the theme of this section. The added term $-\nabla \cdot (|\nabla u|^{p-2} \nabla u)$ is not a mathematical artifice introduced to make a proof work: with $p = 3$ it is the Smagorinsky model, and it is what turbulence modelling has been doing since 1963. The regularisation that makes (A) provable *is* a turbulence model. Fefferman's formulation excludes turbulence from its alternatives, and the price of that exclusion is precisely the term whose absence makes (A) hard — so the difficulty of the Clay problem is manufactured by the same decision that empties its dichotomy.

Turbulence removed, restored, and still not enough. The three observations combine into one. The problem is announced as a problem of turbulence and formulated as a problem without it: the description opens on the wake behind the boat and the air behind the aircraft, and the formal statement that follows asks only whether $|u|$ stays finite. Put the turbulence back in the one form analysis can digest — a turbulent viscosity depending on the local velocity gradients, as in (4) — and Fefferman's (A) ceases to be open. Pointwise solutions exist globally, no singularity forms, and the proof is the standard Sobolev-space argument of [9]: an a priori estimate, a compactness theorem, a monotonicity argument. The turbulence the formalisation discarded is precisely what makes the question answerable.

But it does not reach wellposedness, and this is where the matter rests. What (4) yields is a pointwise solution which exists and is unique. It does not yield a solution whose pointwise values mean anything. A turbulent flow is pointwise unstable whether or not a theorem asserts that its trajectory is unique: perturb the data in the twelfth digit and the velocity at a given point and time is different, while the drag, the lift and the mean pressure are not. Uniqueness in a function space and stability of an output are different properties, and only the second is what a computation, an experiment, or an aircraft depends on. A regularised equation therefore buys the prize's alternative (A) at the cost of answering a question nobody was asking, and leaves untouched the question everybody is: which outputs of this flow are stable, and to what tolerance.

That is the reformulation this article proposes, and the next section states it.

3 What should replace it: wellposedness

The defect in the official formulation is not that it asks a hard question. It is that it asks about the wrong property. Smoothness is not what makes a solution concept physically meaningful. Wellposedness is.

Hadamard's criterion has three parts: a solution exists, it is unique, and it depends continuously on the data. The third is the one that carries the physical content, because data are never exact. A problem whose solution is not stable under perturbation of the data is not a model of anything measurable, whatever its smoothness.

Weak solutions exist. Leray showed in 1934 that weak solutions of (1) exist globally in time [2]. What was left open was not existence but whether the weak solution concept is *wellposed* — whether it determines anything. That is the question worth asking, and it is not the question the Clay problem asks.

Wellposedness is a property of outputs, not of solutions. Here is the point on which the reformulation turns. For a turbulent solution, the pointwise velocity at a given place and time is not determined: perturb the data in the last decimal and the pointwise field is different everywhere within a short time. In that sense the solution is not wellposed and never will be.

But nobody measures the pointwise velocity of a turbulent flow, and no engineering question depends on it. What is measured, and what is asked for, are *mean-value outputs*: the drag on a body, the lift on a wing, a pressure averaged over a surface, a flux through a section. These are functionals of the solution, and the question of wellposedness must be asked of them:

$$M(u) = \int_Q u \cdot \psi \, dx \, dt \tag{6}$$

for a given weight ψ . The claim is that $M(u)$ is stable for turbulent solutions even though u is not.

The mechanism, and why it inverts the classical expectation. Stability of an output is governed by an associated dual, or adjoint, problem: linearise about the computed solution, run it backward from the output functional, and the growth of the dual solution measures how strongly a perturbation of the data or of the residual can move the output. What the dual problem shows is the reverse of what intuition suggests. Over a *turbulent* velocity field — oscillating rapidly in space and time — the dual solution grows only modestly, because the linearised coefficients oscillate and their effect largely cancels. Over a *smooth* field the dual solution grows, because there is nothing to cancel. Turbulence stabilises the outputs that smoothness destabilises.

This is the sense in which the reformulated problem has an answer, and the answer is not the one the dichotomy anticipates. Smooth solutions are not wellposed. Turbulent solutions are wellposed in mean-value outputs. The physically meaningful solution concept is the one the official formulation treats as the failure case.

4 The reformulated problem

We state the replacement plainly.

Problem. Fix smooth divergence-free initial data of size one, and let $\nu > 0$ be arbitrary with no positive lower bound. Show that the weak solution is *wellposed in mean-value outputs, uniformly in ν* : that functionals of the form (6) depend continuously on the data with quantitative error control, by bounds that do not degenerate as $\nu \rightarrow 0$, even though the pointwise solution does not and its gradients diverge as $\nu^{-1/2}$. Equivalently: establish the second law of thermodynamics for fluids — that a turbulent solution dissipates kinetic energy at a rate which does not vanish with the viscosity, and that this dissipation is what makes the outputs stable.

Three features of this statement are worth marking. The *data* are normalised rather than the viscosity, which is the coordinate choice that keeps ν as a parameter and so allows a limit to be taken at all. The bounds are required *uniformly* in ν , which is what makes it a question about turbulence rather than about one Reynolds number at a time. And it asks for stability rather than regularity, so it can be answered affirmatively for the solutions that occur instead of being answered vacuously for solutions that do not.

5 In what sense this constitutes a solution

The solutions are constructed, not assumed. A finite element method for (1) with residual-based stabilisation — Galerkin least squares on a mesh, piecewise polynomial velocity and pressure, no turbulence model supplied by the user (§6) — produces solutions which become turbulent even when initiated from smooth potential data, and which stay bounded. The stabilisation supplies dissipation in proportion to the residual, so the dissipation is generated by the computation rather than prescribed by a model. In the vanishing-viscosity limit the same scheme is an Euler solver whose turbulent dissipation appears automatically. The non-smooth, non-singular third case is therefore not a conjecture: it is what the equations produce when they are solved.

The residual behaves as the theory requires. For such solutions the Navier–Stokes residual $R(U, P)$ is small in a weak norm and large in a strong one. Small weakly is what makes the computed field a solution; large strongly is the statement that it is not smooth. A method whose residual were small in both would be computing something that does not exist.

The error is controlled by duality, not by asserted convergence. An a posteriori error estimate of the form

$$|M(u) - M(U)| \leq \int_Q R(U, P) \cdot \varphi \, dx \, dt \leq S \|R(U, P)\|_{-1}, \quad (7)$$

with φ the dual solution and S its stability factor, converts a computed solution into a statement with a bound attached. The stability factor is not assumed; it is computed by solving the dual

problem. It is through (7) that the wellposedness claim becomes quantitative: S is moderate for mean-value outputs over turbulent solutions and large for pointwise outputs and for smooth solutions, which is the distinction of §3 expressed as a number.

Computation as proof. That an answer rests on computation is not a defect peculiar to this problem. The four-colour theorem was settled the same way, and the settlement was accepted because the computation was reproducible and its error controlled. The same standard applies here, and (7) is what meets it: the claim is not that a particular simulation looked turbulent, but that the output of the computation comes with a bound which can be evaluated, and which is small.

6 The computational record

The argument of §5 is not offered as a programme. The method, the error control and the computations are set out in *Computational Turbulent Incompressible Flow* [5], and it is worth stating what that book establishes, because the reformulation of §4 asks for exactly the things it supplies.

G2: the method is Galerkin, and the turbulence model is the residual. General Galerkin, or G2, is a stabilised finite element method — continuous piecewise linear in space and time, cG(1)cG(1), with least-squares control of the residual. The stabilisation is not a subgrid model: no eddy viscosity is prescribed, no wall function, no closure, no constant fitted to measurement. The dissipation is whatever the residual makes it, which means it is generated by the flow being computed rather than supplied by the person computing it. This matters for the present argument specifically. A computation carrying a fitted turbulence model could not be used to say what Navier–Stokes contains, because the turbulence would have been put in by hand. G2 has nothing put in by hand, so what comes out is a property of (1).

The dual problem is solved, and the stability factor is a computed number. The estimate (7) is the book’s organising device. For each output of interest the dual problem is linearised about the computed solution, solved backwards, and its stability factor evaluated; the mesh is then refined where the product of residual and dual weight is largest. The consequence is that the error in the output is not estimated by comparing successive meshes and hoping, but bounded by a quantity the computation itself produces. Adaptivity and error control are the same mechanism.

Mean-value outputs converge; the flow field does not. The book’s central computational finding is the one §3 needs. Mean drag and mean lift for turbulent flow past a bluff body come out to a tolerance of a few per cent on meshes of a few hundred thousand points — small enough to run on a personal computer, and orders of magnitude coarser than a direct simulation resolving the Kolmogorov scale would require. The pointwise velocity field on those meshes is not converged and does not become so. Both facts hold simultaneously, and together they are the wellposedness claim made quantitative: the outputs are stable because the stability factor for them is moderate, and the pointwise field is not because its stability factor is not.

Turbulent Euler solutions and d’Alembert’s paradox. Taking $\nu \rightarrow 0$ in G2 gives a solver for the Euler equations whose dissipation does not vanish with the viscosity, which is the second law of §4 in computational form. What such solutions produce is not potential flow. A bluff body

in this limit shows potential flow modified by a three-dimensional rotational separation, and that separation carries drag — which is the resolution of d’Alembert’s paradox offered in [6]: the paradox arises from taking the smooth potential solution to be the vanishing-viscosity limit, and the correct limit is turbulent. The same analysis applied to a wing gives lift and drag without a circulation condition imposed by hand [7].

This is the concrete content behind the claim of §2 that smooth solutions are unstable. It is not an assertion about what would happen to a perturbed smooth solution; it is a record of what happens when the equations are solved from smooth data without a model to keep them smooth.

7 What the answer says about the original question

To the official question — do smooth data give solutions that remain smooth for all time — the answer implied by the above is no, and not by blowup. Smooth data give solutions that become turbulent: non-smooth, with gradients growing as $\nu^{-1/2}$, bounded in velocity, dissipating at a rate which does not vanish as $\nu \rightarrow 0$. Blowup, the only failure mode the official formulation admits, does not occur and is not the interesting question. The interesting question is the one the formulation leaves out.

This does not claim the prize, and it is not intended to. It claims something less procedural and more useful: that the problem as posed cannot be answered in a way that illuminates fluid mechanics, because the alternative it offers is not exhaustive over the cases that occur; and that the adjacent, answerable question — wellposedness of mean-value outputs of turbulent weak solutions — is both the physically meaningful one and one to which the answer is affirmative and computable.

References

- [1] C. Fefferman, *Existence and smoothness of the Navier–Stokes equation*, Official Problem Description, Clay Mathematics Institute, 2000.
- [2] J. Leray, *Sur le mouvement d’un liquide visqueux emplissant l’espace*, Acta Math. **63** (1934) 193–248.
- [3] P. Constantin, W. E and E. S. Titi, *Onsager’s conjecture on the energy conservation for solutions of Euler’s equation*, Comm. Math. Phys. **165** (1994) 207–209.
- [4] P. Isett, *A proof of Onsager’s conjecture*, Ann. of Math. **188** (2018) 871–963.
- [5] J. Hoffman and C. Johnson, *Computational Turbulent Incompressible Flow*, Applied Mathematics: Body and Soul, vol. 4, Springer, 2007.
- [6] J. Hoffman and C. Johnson, *Resolution of d’Alembert’s paradox*, J. Math. Fluid Mech. **12** (2010) 321–334.
- [7] J. Hoffman, J. Jansson and C. Johnson, *New theory of flight*, J. Math. Fluid Mech. **18** (2016) 219–241.
- [8] J. Hoffman, *Efficient computation of mean drag for the subcritical flow past a circular cylinder using general Galerkin G2*, Int. J. Numer. Meth. Fluids **59** (2009) 1241–1258.
- [9] J.-L. Lions, *Quelques méthodes de résolution des problèmes aux limites non linéaires*, Dunod, Paris, 1969.

- [10] C. Johnson, posts under the label *clay problem*, <https://claesjohnson.blogspot.com/search/label/clay\%20problem>.